

Request Flow Summary

End-to-End Trace Through the System

Generated: February 08, 2026

1. Complete Request Flow

This diagram shows the complete journey of a request through the system, from HTTP input to streaming output:

Step	Component	Description
1	HTTP Request	FastAPI receives incoming request
2	Frontend	Tokenizes text into input_ids
3	Backend	Receives UserMsg with input_ids
4	PrefillManager	Adds request to pending_list queue
5	Scheduler	Creates Batch with Req objects
6	Engine	Executes forward pass (prefill/decode)
7	Sampler	Generates next_tokens from logits
8	Detokenizer	Converts tokens back to text
9	Frontend	Creates UserReply with incremental_output
10	HTTP Response	Streams response back to client

2. Key Object Transformations

The request object transforms through multiple types as it flows through the system:

```
GenerateRequest (HTTP)
↓
TokenizeMsg (frontend)
↓
UserMsg (backend, has input_ids)
↓
PendingReq (waiting in queue)
↓
Req/ChunkedReq (scheduled for execution)
↓
Batch (grouped for GPU)
```

```

↓
ForwardOutput (next tokens)
↓
DetokenizeMsg (token to text)
↓
UserReply (final response)

```

Object Type	Key Fields	Purpose
GenerateRequest	prompt, max_tokens, temperature	HTTP API request
TokenizeMsg	uid, text, sampling_params	Tokenization request
UserMsg	uid, input_ids, sampling_params	Backend processing
PendingReq	uid, input_ids, chunked_req	Queued request
Req	input_ids, output_len, cache_handle	Scheduled request
Batch	reqs[], phase, input_ids	GPU batch processing
ForwardOutput	next_tokens_gpu, next_tokens_cpu	Generated tokens
DetokenizeMsg	uid, next_token, finished	Token to detokenize
UserReply	uid, incremental_output, finished	Streaming response

3. Request State Transitions

Each request goes through distinct states during its lifecycle:

```
WAITING: PendingReq in pending_list  
↓  
PREFILL: Req in Batch(phase="prefill")  
↓  
DECODE: Req in Batch(phase="decode")  
↓  
FINISHED: finished=True in DetokenizeMsg/UserReply
```

State	Location	Description
WAITING	PrefillManager.pending_list	Request queued, waiting for GPU resources
PREFILL	Batch with phase='prefill'	Computing KV cache for input tokens
DECODE	Batch with phase='decode'	Generating output tokens one by one
FINISHED	UserReply.finished=True	Generation complete, response sent

4. The 'yield' Pattern in Streaming

The system uses Python generators with 'yield' to enable token-by-token streaming. This is why you see tokens appear one-by-one instead of waiting for the complete response.

wait_for_ack function:

```
# In wait_for_ack:  
for ack in pending:  
    yield ack # Returns UserReply one by one  
  
# Each yield pauses execution and returns one item  
# Next call resumes from where it left off
```

stream_generate function:

```
# In stream_generate:  
async for ack in self.wait_for_ack(uid):  
    yield f"data: {ack.incremental_output}\n"  
    # Each token streamed immediately!  
    if ack.finished:  
        break
```

Aspect	Without yield	With yield
Return behavior	Returns all at once	Returns one item at a time
Memory usage	Holds entire result in memory	Processes items on-demand
User experience	Wait for complete response	See tokens as they generate

Network traffic	Single large payload	Multiple small chunks (SSE)
-----------------	----------------------	-----------------------------

5. Key Request Object Fields

Important fields tracked throughout the request lifecycle:

Field	Type	Description
input_ids	torch.Tensor	Tokenized input (e.g., [101, 2023, 2003, ...])
output_ids	torch.Tensor	Generated tokens appended to input (grows during decode)
num_computed_tokens	int	Tracks KV cache length (how many tokens cached)
sampling_params	SamplingParams	Controls temperature, top_p, max_tokens, etc.
cached_len	int	Number of tokens with computed KV cache
output_len	int	Number of output tokens generated so far
table_idx	int	Index in attention table for this request
cache_handle	BaseCacheHandle	Handle to KV cache storage

6. Complete Code Flow Example

Here's how the key functions connect in the actual codebase:

```
# 1. HTTP endpoint receives request
@app.post("/generate")
async def generate(req: GenerateRequest):
    uid = state.add_request(req)
    return StreamingResponse(state.stream_generate(uid))

# 2. Frontend tokenizes and sends to backend
tokenize_msg = TokenizeMsg(uid=uid, text=req.prompt, ...)
backend_msg = UserMsg(uid=uid, input_ids=tokens, ...)

# 3. Backend adds to pending queue
def _process_one_msg(self, msg: UserMsg):
    self.prefill_manager.add_one_req(msg)

# 4. Scheduler creates batch
batch = self.prefill_manager.schedule_next_batch()

# 5. Engine forward pass
forward_output = self.engine.forward_batch(batch)

# 6. Detokenize and reply
detokenize_msg = DetokenizeMsg(uid=uid, next_token=token)
user_reply = UserReply(uid=uid, incremental_output=text)

# 7. Stream response
async for ack in self.wait_for_ack(uid):
    yield f"data: {ack.incremental_output}\n"
```

Summary

The request flow demonstrates a well-architected pipeline that efficiently handles LLM inference with streaming responses. Key design patterns include:

- **Asynchronous Processing:** Non-blocking I/O for concurrent request handling
- **Batch Processing:** Grouping requests for efficient GPU utilization
- **Streaming Output:** Using generators (`yield`) for real-time token delivery
- **State Management:** Clear state transitions (`WAITING` → `PREFILL` → `DECODE` → `FINISHED`)
- **Type Safety:** Strongly typed messages at each pipeline stage
- **Resource Management:** KV cache handles and table management for memory efficiency